

CORTEX

A LITTLE MORE HUMAN.

A robot that hears
how you speak.

The same words.


A different voice.

A different physical response.



Same instruction. Different execution.

CALM VOCAL PROFILE




“Come here.”

0.55 m/s ceiling

Normal priority · measured approach · 1.2 m space

URGENT VOCAL PROFILE



“Come here.”

0.90 m/s ceiling

Higher priority · faster approach · shorter response

ACTUAL PHYSICS TEST / 4 SIMULATED SECONDS: 1.88 m CALM → 2.98 m URGENT

Reason slowly. Stop immediately.

CORTEX / ARCHITECTURE

Voice becomes physical policy.

The reflex can stop the robot while the planner is still thinking.



Execution: local / EdgeOne adapter

Trace every decision. Measure every latency.

Voice never overrides safety.

Structured tools. Bounded actions. Stale commands expire. No model sends motor torques.

Command acknowledgement and physical settling are measured separately.

“Where did you last see my backpack?”

A remembered frame becomes a place the robot can return to.

- 01 Capture and index a scene frame
- 02 Retrieve the original sighting
- 03 Navigate to its registered waypoint



BACKPACK → PLANTER

FRAME EVIDENCE, NOT A GUESS

“I saw it beside the planter.”

Then: “Take me there.”

Physical AI should hear the human behind the instruction.

Built and exercised

- ↗ Actual G1 walking, stopping and turning
- ↗ Acoustic-conditioned speed and approach
- ↗ Evidence retrieval, confirmation and reflexes
- ↗ Inspect every policy, trace and measured delay

Ready for the next validation

- ↗ Live same-words microphone trial with Gadium
- ↗ Yerba Buena smartphone capture → neural scene
- ↗ Credentialed SambaNova / Memories.ai runs
- ↗ Photorealism review and remote runtime setup

CORTEX

Hear. Remember. Reason. React. Safely act.